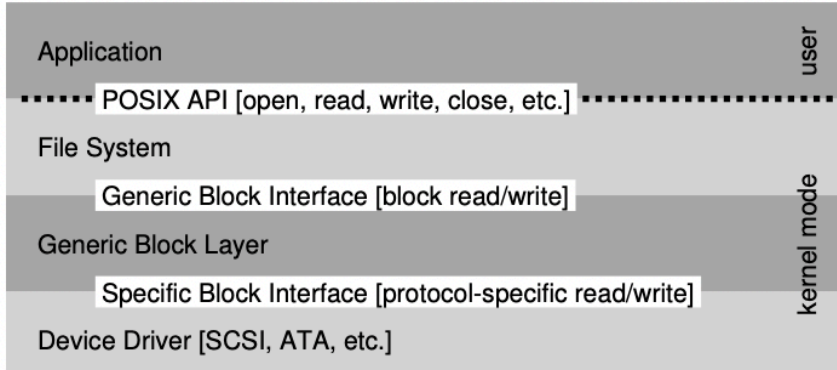


CS111	AD
Chapter 2	Intro to OS
Von Neuman model	Von Neuman model, the processor fetches an instruction, decodes it, and then executes the instruction.
Resource Manager	An OS is in charge of making sure the system operates correctly and efficiently in an easy to use manner.
Virtualizing the CPU	The illusion that a system has many virtual CPU's can turn a small number of physical CPU's into a seemingly infinite number of CPU's which then allows a programs to seemingly run at once is called virtualizing the CPU
Virtualization	Each process accesses its own private virtual address space which the OS maps onto the physical memory of the machine The OS manages memory such that one running program does not affect the memory address space of other processes. The program itself thinks that it has all the memory to itself.
Concurrency	A broad term to refer to a host of problems that arise when working on multiple problems or processes at once. In an example correlating to increment a counter by using two threads, for very high values of the counter, we see that the program does not work as expected. This is because the load, increment, and store instructions do not happen <u>automatically</u> (not all at once) the program does not work. Concurrency deals with such issues.
Persistence	Persistence is the idea that even when power goes away we need hardware and software to be able to store data despite this.
File System	The part of the OS that manage the disk is the file system, and it is responsible for storing any files that the user creates in a reliable and efficient manner.
Device Driver	A device driver is a piece of code in the OS that can deal with a specific device.
Journaling, Copy-on-write	The OS tends to bunch up write instructions for files since it becomes more effective, however to better deal with crashes there are intricate protocols called journaling or copy-on-write.

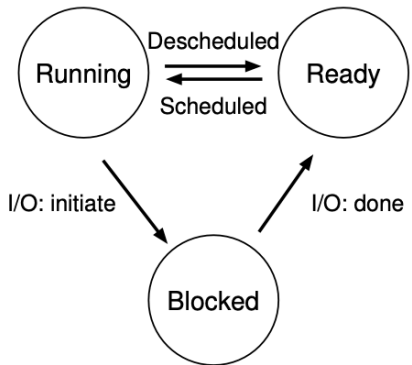
Goals of the OS	<p>The main goals of OS systems are as follows:</p> <ul style="list-style-type: none"> - Abstraction: make the system easy to use - Performance: minimize the overheads - Protection: between applications, OS and applications. - Isolation: isolating processes from one another is key to protection - Energy-Efficiency - Security - Mobility
History of OS	<p>Initially, the OS was just a set of libraries of commonly used functions used by different programmers. It used a mode called <u>batch</u> processing where a number of jobs were set up and then run in bulk by an operator.</p> <p>Then protection was introduced. It was done by introducing the hardware <u>privilege level</u>, which means the hardware restricts can do, for example a user mode function cannot initiate I/O requests to the disk.</p> <p>When an I/O request is made, an instruction in the hardware called a <u>trap</u> which would raise the the privilege level to <u>kernel mode</u>. In kernel mode, the OS has full access to the hardware of the system. Then a <u>return-from-trap</u> instruction reverts to user <u>mode mode</u>.</p> <p>Then the idea of multiprogramming became more prevalent, which is where issues like concurrency and memory protection became important</p>

Chapter 36	I/O Devices
<p>System Architecture</p> <ul style="list-style-type: none"> - Lower performance devices components are placed farther away eg. (IO devices etc) - Higher performance devices are generally placed closer to the CPU eg. graphics, memory etc. 	<pre> graph TD CPU[CPU] --- MB[Memory Bus] Memory[Memory] --- MB MB --- GIB[General I/O Bus] GIB --- Graphics[Graphics] GIB --- PIOB[Peripheral I/O Bus] PIOB --- D1[(Disk 1)] PIOB --- D2[(Disk 2)] PIOB --- D3[(Disk 3)] PIOB --- D4[(Disk 4)] </pre> <p>The diagram illustrates a hierarchical system architecture. At the top, the CPU and Memory are connected to a horizontal line representing the Memory Bus (proprietary). This bus then branches down to a General I/O Bus (e.g., PCI). The General I/O Bus is connected to the Graphics component. Below the General I/O Bus is the Peripheral I/O Bus (e.g., SCSI, SATA, USB), which is connected to four disk drives represented by cylinder icons.</p>

Canonical Device	<p>Hardware interface is something that allows the system's software to control its operation. The second part of any device is its internal structure, which implements the abstraction the device presents to the system.</p> <p>Firmware is the software within a hardware device that implements it's functionality</p> <p>Polling a device essentially means, repeatedly reading the status register, to understand what the state of the machine is. (Programmed I/O PIO)</p>																																																												
How can the OS check device status without frequent polling?	<p>When an I/O call is made, the CPU may have to sit idle while not doing anything. To improve this we could use an <u>interrupt handler</u> and interrupt service routine, which allows for overlap with computation and I/O</p> <div><div>CPU</div><table><tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>p</td><td>p</td><td>p</td><td>p</td><td>p</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table><div>Disk</div><table><tr><td></td><td></td><td></td><td></td><td></td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td></td><td></td><td></td><td></td><td></td></tr></table></div> <div><div>CPU</div><table><tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>2</td><td>2</td><td>2</td><td>2</td><td>2</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table><div>Disk</div><table><tr><td></td><td></td><td></td><td></td><td></td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td></td><td></td><td></td><td></td><td></td></tr></table></div>	1	1	1	1	1	p	p	p	p	p	1	1	1	1	1						1	1	1	1	1						1	1	1	1	1	2	2	2	2	2	1	1	1	1	1						1	1	1	1	1					
1	1	1	1	1	p	p	p	p	p	1	1	1	1	1																																															
					1	1	1	1	1																																																				
1	1	1	1	1	2	2	2	2	2	1	1	1	1	1																																															
					1	1	1	1	1																																																				
Balancing Interrupts, and PIO	<p>Suppose that a device is very fast then what results is that the cost of receiving of an interrupt, handling it, and then switching back to the issuing process may actually end up taking up more time</p> <p>Therefore depending on the device, a hybrid that polls for a little while, and if the device is not yet finished, uses interrupts.(Two phase approach)</p>																																																												
Livelock	<p>When a huge stream of incoming packets each generate an interrupt, it is possible for the OS to only end up processing interrupts and then never allowing the user-level processes to run.</p>																																																												
Coalescing	<p>In this setup a device waits before raising an interrupt, which while it waits allows other requests to complete, and allows for multiple interrupts to be merged into one.</p>																																																												
DMA	<p>Direct Memory Access (DMA) is a specific device that orchestrates transfers between devices and memory without CPU intervention. The OS would first tell the DMA where the data is in memory, and then how much to copy, and which device to send it to. When the DAM is complete the DMA raises an interrupt indicated to the OS that the transfer is complete</p> <div><div>CPU</div><table><tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>2</td><td>2</td><td>2</td><td>2</td><td>2</td><td>2</td><td>2</td><td>2</td><td>1</td><td>1</td></tr></table><div>DMA</div><table><tr><td></td><td></td><td></td><td></td><td></td><td>c</td><td>c</td><td>c</td><td></td><td></td><td></td><td></td><td></td><td></td><td></td></tr></table><div>Disk</div><table><tr><td></td><td></td><td></td><td></td><td></td><td></td><td></td><td></td><td></td><td></td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table></div>	1	1	1	1	1	2	2	2	2	2	2	2	2	1	1						c	c	c																		1	1	1	1	1															
1	1	1	1	1	2	2	2	2	2	2	2	2	1	1																																															
					c	c	c																																																						
										1	1	1	1	1																																															

How should hardware communicate with a device?	<p>First method is to have explicit I/O instructions; these instructions specify a way for the OS to send data to a specific device. Such instructions are considered privileged, meaning only the OS can use them.</p> <p>The second method is memory mapped I/O. The hardware makes device registers as though they were actually memory locations. So instead the OS issues a load or store instruction to that specific address.</p>
Making a device neutral OS	 <p>The diagram illustrates a layered architecture for a device-neutral OS. It is divided into two main sections: 'user' (top) and 'kernel mode' (bottom). The 'user' section contains the 'Application' layer, which interacts with the 'File System' layer via the 'POSIX API [open, read, write, close, etc.]'. The 'kernel mode' section contains the 'Generic Block Layer', which interacts with the 'Device Driver [SCSI, ATA, etc.]' via the 'Specific Block Interface [protocol-specific read/write]'. The 'File System' layer interacts with the 'Generic Block Layer' via the 'Generic Block Interface [block read/write]'.</p>
Notes for 36.8 and 36.9 not included.	<p>The issue of making a device neutral OS is that device drivers with specialized features may end up being unused since the kernel and OS would require a generalized interface from the driver</p>

Chapter 4	Abstraction of Processes
Process	A process is just the running of a program, the OS manages to execute the bytes and get them running.
Virtualization, and time sharing	Gives the illusion of having many CPUs when in reality there is only one CPU. Through time sharing of the CPU, the OS allows users to run as many concurrent processes as they would like ⇒ however the cost is performance since as there are more concurrent processes, the more slowly the CPU will be shared
Space Sharing	Where a resource is divided (in space) among those who use it. For example a file takes up part of disk storage which is a naturally space shared resource
Context Switch	This enables the OS to stop running one program and start running another different program on a given CPU.
Mechanism Vs Policy	Mechanism deals with the “how”, policy deals with “Which” question to answer. Having this difference is a form of modularity that allows policies to change without a changing mechanism.

<p>Machine State of a process:</p>	<p>The machine state is what a program can read or update when it is running.</p> <ul style="list-style-type: none"> • Memory: the address space of a process, that part of memory that a program can read and write to. • The instruction pointer, stack pointer, and frame pointer, help tell which instruction should be executed next. And manage the environment the code should get executed in (local variables, function calls etc). • I/O information files, or external hardware devices
<p>Process API</p>	<p>Create: An OS must have some method to create new processes. Destroy: halt or kill a processes that is runaway (some way to destroy a process forcefully) Wait: Waiting for a process to finish it's execution Miscellaneous Control: Give the ability to pause a process for sometime and then resume running at a later time Status: get information about the process. Ex how long it has run for, or what state it is in.</p>
<p>Process Creation:</p>	<p>The OS will first load code into the address space from the disk so that process can execute its instructions.</p> <ul style="list-style-type: none"> • Eager loading: all loading is done at once before running the program • Lazily: loaded as and when a program needs, uses paging and swapping to properly work <p>The OS will then load and allocate memory for the program's run-time stack, the program's heap.</p> <ul style="list-style-type: none"> ❖ The OS also will have the ability to give more memory to heap, if a program calls malloc() <p>Finally the OS transfers control the program, by jumping to the main() routine.</p>
<p>Process States:</p>	<p>Running: The processing is executing instructions Ready: the process is ready to run, but the OS is not letting it execute any instructions. Blocked: the process is not ready to run, until some other event take places</p>  <pre> graph TD Running((Running)) -- "Descheduled" --> Ready((Ready)) Ready -- "Scheduled" --> Running Running -- "I/O: initiate" --> Blocked((Blocked)) Blocked -- "I/O: done" --> Ready </pre> <p>The diagram illustrates the transitions between process states. It features three circles labeled 'Running', 'Ready', and 'Blocked'. A double-headed arrow connects 'Running' and 'Ready', with 'Descheduled' above and 'Scheduled' below. An arrow points from 'Running' to 'Blocked' labeled 'I/O: initiate'. Another arrow points from 'Blocked' to 'Ready' labeled 'I/O: done'.</p>

How does the OS keep track of processes?	<p>The OS has multiple data structures that would allow it to store information about each process.</p> <pre> // the registers xv6 will save and restore // to stop and subsequently restart a process struct context { int eip; int esp; int ebx; int ecx; int edx; int esi; int edi; int ebp; }; // the different states a process can be in enum proc_state { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE }; // the information xv6 tracks about each process // including its register context and state struct proc { char *mem; // Start of process memory uint sz; // Size of process memory char *kstack; // Bottom of kernel stack // for this process enum proc_state state; // Process state int pid; // Process ID struct proc *parent; // Parent process void *chan; // If non-zero, sleeping on chan int killed; // If non-zero, have been killed struct file *ofile[NOFILE]; // Open files struct inode *cwd; // Current directory struct context context; // Switch here to run process struct trapframe *tf; // Trap frame for the // current interrupt }; </pre> <p>In this example, we also see the context, where for each process the OS would store the contents of the registers if a process needed to be stopped, and started at a later point ⇒ this helps enable context switching</p> <p>The OS will contain a process list (a list of all the processes) and each process will have a Process Control Block (PCB) which is a C structure that contains information about each process.</p>
--	--

Chapter 6	Limited Direction Execution
Limited Direction Execution	<p>Direct Execution implies that the program runs natively on the CPU, and then transfers control back to the OS. However this brings up two issues, how do we stop the program from doing things we don't want and secondarily how do we implement time sharing.</p>

Problem #1 Creating Restricted Operations	When running a program, we may not want to limit its access for security and efficiency, so how can we do this?																					
Protected Control Transfer:	The hardware helps create different modes (kernel and user mode) and provides trap, and return from trap instructions to change between the kernel and the user mode programs.																					
System Call	When a user wants to execute an instruction that is in kernel mode (such as I/O) , they make a system call																					
Traps/ and Switching	To execute a system call, a <u>trap</u> instruction is executed, which jumps into the kernel and raises the privilege level. To store the user program states, the counter, flags and other information will be pushed onto the kernel stack																					
Return from Trap	After the kernel has finished running the special instructions, the return from trap pops the above values off the stack which resumes execution of the user systems.																					
Trap Table	To prevent the user from executing its own code instead of what the trap should do, the kernel sets up a trap table. The trap table tells what code the hardware should execute when certain trap events occur. The kernel then gives the location of these trap handlers to the hardware.																					
Problem #2 How the OS regains control	A cooperative process: whenever a process calls a system call, (which contains a yield ⇒ specifically for just transferring control) or if a process does something illegal, control will be automatically transferred ex divide by zero																					
Timer interrupt	<p>An uncooperative process: uses a timer interrupt, every few milliseconds, an interrupt is raised, and a process is halted and the OS runs. (However this much more of a hardware feature not a software one) the OS cannot do anything from a software side without this hardware implementation.</p> <table><tr><td>OS @ boot (kernel mode)</td><td>Hardware</td><td></td></tr><tr><td>initialize trap table</td><td>remember addresses of... syscall handler timer handler</td><td></td></tr><tr><td>start interrupt timer</td><td>start timer interrupt CPU in X ms</td><td></td></tr><tr><td>OS @ run (kernel mode)</td><td>Hardware</td><td>Program (user mode)</td></tr><tr><td></td><td></td><td>Process A ...</td></tr><tr><td></td><td>timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler</td><td></td></tr><tr><td>Handle the trap Call switch() routine save regs(A) to proc-struct(A) restore regs(B) from proc-struct(B) switch to k-stack(B) return-from-trap (into B)</td><td>restore regs(B) from k-stack(B) move to user mode jump to B's PC</td><td>Process B ...</td></tr></table>	OS @ boot (kernel mode)	Hardware		initialize trap table	remember addresses of... syscall handler timer handler		start interrupt timer	start timer interrupt CPU in X ms		OS @ run (kernel mode)	Hardware	Program (user mode)			Process A ...		timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler		Handle the trap Call switch() routine save regs(A) to proc-struct(A) restore regs(B) from proc-struct(B) switch to k-stack(B) return-from-trap (into B)	restore regs(B) from k-stack(B) move to user mode jump to B's PC	Process B ...
OS @ boot (kernel mode)	Hardware																					
initialize trap table	remember addresses of... syscall handler timer handler																					
start interrupt timer	start timer interrupt CPU in X ms																					
OS @ run (kernel mode)	Hardware	Program (user mode)																				
		Process A ...																				
	timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler																					
Handle the trap Call switch() routine save regs(A) to proc-struct(A) restore regs(B) from proc-struct(B) switch to k-stack(B) return-from-trap (into B)	restore regs(B) from k-stack(B) move to user mode jump to B's PC	Process B ...																				

Chapter 7	Scheduling: Introduction
<p>Workload:</p> <p>Turnaround time</p> <p>First In, First Out (FIFO)</p> <p>Shortest Job First:</p> <p>Shortest Time-to-Completion First:</p> <p>Round Robin (RR):</p> <p>Section 7.9 TBD</p>	<p>Simplifying Assumptions about the process running in the system</p> <p>The time at which the job completes minus the time at which the job arrived: $\text{Turnaround} = T_{\text{completion}} - T_{\text{arrival}}$</p> <p>The first task that comes in, will be executed until it is completed and so on and so on.</p> <p>Disadvantage: convoy effect, where a number of much shorter potential processes end up getting queued behind a process that will run for much longer.</p> <p>The shortest job runs first, this works well under the assumption that all processes enter at the same time.</p> <p>Disadvantage: convoy effect, when processes don't arrive at the same time.</p> <p>Every time a new process enters the queue it determines which task has the smallest time to completion, and then schedules that.</p> <p>Disadvantage: horrendous response time and very bad for interactivity, when people want to see their results as soon as possible.</p> <p>RR runs each job for a time slice (a quantum) and then switches to the next job in the run queue. However an important note is that the time slice must be a multiple of the timer-interrupt period. RR is a fair policy since it evenly splits up the CPU</p> <p>Disadvantage: RR is horrendous under the metric of turnaround time</p>

Chapter 27	Concurrency: Thread API
Thread creation	<p>Int pthread_create(pthread_t * thread, attr, void * (*start_routine) (void*), void * args);</p> <p>The first argument, is the thread data structure to store information about the thread, the second argument are about the attributes of the thread, the *star routine, is a function pointer to the function that the thread should start executing, and the final parameter is the arguments given to a thread.</p>
Waiting for a thread to complete:	<p>Int pthread_join(pthread * thread, void* return), one argument to specify which thread to wait for, and a second argument to see the return value.</p>
Return values from a thread:	<p>We cannot return a pointer to something that has been allocated on the threads stack, since that would end up returning a pointer to a memory location that is now being used for something else once the thread finishes executing.</p>
Mutex Locks	<p>We can add a mutex lock around a critical section to indicate that only one thread can execute it at a time.</p> <pre>pthread_mutex_t lock; Int rc = pthread_mutex_init(&lock, NULL); pthread_mutex_lock(&lock); assert(rc == 0); //If success x = x + 1 //Critical Section here Ptheard_mutex_unlock(&lock)</pre>
Condition Variables	<p>Condition variables help with signaling between multiple threads, we can have one thread be “sleeping” until a certain condition is met, and when another thread changes the condition, that thread can wake up this other thread.</p> <p>We must hold the lock, until after we wake up the thread in-order to prevent any race conditions.</p> <p>To execute a program with multiple threads under gcc we use the #include<pthreads.h> header</p> <p>Sleeping Process:</p> <pre>pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER; Pthread_cond_t init = PTHREAD_COND_INITIALIZER; //Lock the lock while(initialized == 0) pthread_cond_wait(&init, &lock)</pre> <p>Waking Process:</p> <pre>pthread_cond_signal(&init)</pre>

Chapter 28	Concurrency: Locks
Locks	<p>The lock() attempts to acquire the lock, if there is somebody else holding onto the lock, then this attempt will fail.</p> <p>The unlock() will release the lock, so that somebody else may acquire the lock.</p>
Coarse versus Fine grained	<p>In a coarse-grained approach we will protect a critical section by using one large lock, in a fine-grained approach we will have different locks protecting different data and different data structures, thereby increasing concurrency.</p>
Turning off Interrupts	<p>The first approach for mutual exclusion was to disable interrupts for critical sections, thereby ensuring that the code will be executed fully.</p> <p>Negatives: we have to give a program a privileged operation, which can allow a program to monopolize a system and thereby run forever. Lastly, disabling and enabling can significantly slow down a system.</p>
Test-and-set	<p>An atomic instruction that enables us to test a value by reutring it, while also changing it to a different new value</p>
Spin Lock:	<p>Uses the test and set function to implement a lock, and will have a thread just spin until it can acquire the lock.</p> <p>However spin locks are not fair, and can lead to starvation since a waiting thread may never get the lock</p>
Compare and Swap:	<p>Similar to a spin lock, however, instead of always setting the value to 1, we check whether the lock is free and then set the value to 1.</p>
Load and Store	<p>Load the value and check if it is free, if so, attempt to store the value with a conditional check to see if it has already been acquired.</p>
Fetch and Add, Ticket Lock	<p>Helps implement a ticket lock, whenever a thread wants to access a lock it will increment a global counter, and store it as it's turn, when it reaches that number, it gets to use the resource.</p> <p>Ticket locks, still have spinning, but they prevent starvation.</p>
Solution#1 to Spinning: Yield	<p>Instead of having a process just constantly be in a while loop spinning, within the while loop we will yield the CPU to a different process. While this does help reduce wasted time, it also requires a lot of context switching.</p>
Lock with Queues:	<p>We will use a queue, to hold onto a lock, and this way we have control over who next gets the lock</p>
Two Phase Locks:	<p>First phase, spin for a bit, second phase yield, and put caller to sleep</p>

Chapter 31	Concurrency: Semaphores
Initializing A Semaphore	<pre data-bbox="560 226 982 336">#include <semaphore.h> sem_t s; sem_init(&s, 0, 1);</pre> <p data-bbox="560 378 1437 487">The third argument will initialize the value of the semaphore to that value, and the second argument value of 0 is used to indicate that the semaphore is shared between threads, in the same process.</p>
Value of a Semaphore	<p data-bbox="560 531 1502 598">A negative value for a semaphore indicates the number of threads that are waiting.</p> <pre data-bbox="560 609 1494 745">int sme_wait(sem_t *s) { Decrement the value of semaphore s by one Wait if the value of semaphore s is negative }</pre> <pre data-bbox="560 798 1437 934">int sme_post(sem_t *s) { Increment the value of semaphore s by one If there are more thread waiting wake one }</pre>
Binary Semaphore (Lock)	<p data-bbox="560 951 1485 1060">We give our semaphore an initial value of 1 to make in a binary semaphore. To try and acquire the lock we will use <code>sem_wait()</code>, and then to release the lock we will use <code>sem_post()</code></p>
Condition Variable	<p data-bbox="560 1098 1469 1207">We give our semaphore an initial value of 0. The parent (or the waiting process) will call <code>sem_wait()</code>, while the child process or thread (the signaling process) will call <code>sem_post()</code></p>
Bounded Buffer Problem	<pre data-bbox="560 1239 1437 1963">1 sem_t empty; 2 sem_t full; 3 sem_t mutex; 4 5 void *producer(void *arg) { 6 int i; 7 for (i = 0; i < loops; i++) { 8 sem_wait(&empty); // line p1 9 sem_wait(&mutex); // line p1.5 (MOVED MUTEX HERE...) 10 put(i); // line p2 11 sem_post(&mutex); // line p2.5 (... AND HERE) 12 sem_post(&full); // line p3 13 } 14 } 15 16 void *consumer(void *arg) { 17 int i; 18 for (i = 0; i < loops; i++) { 19 sem_wait(&full); // line c1 20 sem_wait(&mutex); // line c1.5 (MOVED MUTEX HERE...) 21 int tmp = get(); // line c2 22 sem_post(&mutex); // line c2.5 (... AND HERE) 23 sem_post(&empty); // line c3 24 printf("%d\n", tmp); 25 } 26 } 27 28 int main(int argc, char *argv[]) { 29 // ... 30 sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with... 31 sem_init(&full, 0, 0); // ... and 0 are full 32 sem_init(&mutex, 0, 1); // mutex=1 because it is a lock 33 // ... 34 }</pre>

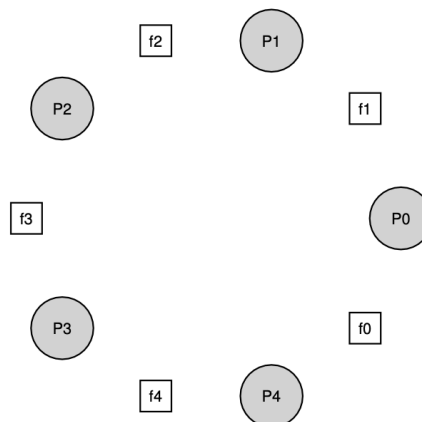
Read Write locks

Allow as many readers to read the data without modification without letting any writers in. Then when a writer enters, there can be only one person writing to a data structure at a time.

```
1  typedef struct _rwlock_t {
2      sem_t lock;          // binary semaphore (basic lock)
3      sem_t writelock;    // used to allow ONE writer or MANY readers
4      int  readers;       // count of readers reading in critical section
5  } rwlock_t;
6
7  void rwlock_init(rwlock_t *rw) {
8      rw->readers = 0;
9      sem_init(&rw->lock, 0, 1);
10     sem_init(&rw->writelock, 0, 1);
11 }
12
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock); // first reader acquires writelock
18     sem_post(&rw->lock);
19 }
20
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock); // last reader releases writelock
26     sem_post(&rw->lock);
27 }
28
29 void rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32
33 void rwlock_release_writelock(rwlock_t *rw) {
34     sem_post(&rw->writelock);
35 }
```

Dining Philosophers

If each philosopher were to grab the fork on the left, none of them would have the ability to eat since they would each have exactly one fork. TO AVOID THIS DEADLOCK WE SIMPLY NEED TO CHANGE ONE PHILOSOPHER to try and grab a fork on the right first.



Implementing A Zempahore

```
1  typedef struct __Zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  } Zem_t;
6
7  // only one thread can call this
8  void Zem_init(Zem_t *s, int value) {
9      s->value = value;
10     Cond_init(&s->cond);
11     Mutex_init(&s->lock);
12 }
13
14 void Zem_wait(Zem_t *s) {
15     Mutex_lock(&s->lock);
16     while (s->value <= 0)
17         Cond_wait(&s->cond, &s->lock);
18     s->value--;
19     Mutex_unlock(&s->lock);
20 }
21
22 void Zem_post(Zem_t *s) {
23     Mutex_lock(&s->lock);
24     s->value++;
25     Cond_signal(&s->cond);
26     Mutex_unlock(&s->lock);
27 }
```

A Zemaphore is almost exactly same to a semaphore, however a zemapahore does not maintain the invariant that if the value is negative it will represent the number of waiting processes. In this case it will always be positive

Signals

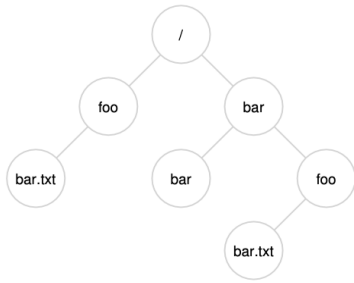
```
#include <signal.h>
void handle(int arg) {
    printf("Handle the signal")
}

int main(){
    signal(SIGHUP, handle); //set up the signal
    ...
}
```

Whenever our program catches a signal it will stop, call the handle function, and then once again resume execution.

Chapter 32	Concurrency: Common Concurrency Problems
Non Deadlock bugs	Bugs in concurrency programs that don't actually have the system end up in deadlock.
Atomicity Violation	When we memory access and modify a shared location, atomicity is not enforced. This can be fixed by adding locks around the respective critical sections.
Order Violation Bugs	The memory access order of two groups is flipped, this can be easily fixed by adding a conditional variable to indicate when the status has changed.
Conditions for Deadlock (all four must occur)	<p>Mutual Exclusion: Threads claim exclusive control of resources that they require</p> <p>Hold and Wait: Threads hold resources allocated to them, while waiting for additional resources</p> <p>No preemption: Resources cannot be forcibly removed from threads that are holding them.</p> <p>Circular Wait: There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain.</p>
Solving Deadlock:	<p>Avoid circular wait: by making sure that locks can only be acquired in a very certain order.</p> <p>Avoid hold and wait: acquire all locks at once atomically.</p> <p>No preemption: we can use a trylock, which will tell use if a lock is available, and if it not release our current resource as well, however this can lead to an issue of livelock</p> <p>The final option is to avoid deadlock via scheduling, however this would require the scheduler to know what threads would need which resources.</p>

Chapter 39	File and Directories
File	<p>A file is simply a linear array of bytes, each of which has the read and write permissions.</p> <p>Each file has a "name" which is called the inode number, which is a number to identify the file that the user cannot see</p>
Directory	A directory also has an inode number, and its contents are a list of (file_name, inode number). The file_name is what the user can see.



We can place directories within directories to create a Directory Tree or a Directory Hierarchy

Directories and files can have the same name as long as they are in different locations.

Directories start at the root directory: “/”. Each file has an absolute pathname.

Creating Files:

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);
```

open returns a file descriptor, if the file foo does not exist, it will create it with the flag O_CREAT, and the O_WRONLY can only be written too. If the file already exists the O_TRUNC states to set the newly created file to have zero bytes

Reading a file:

```
open("foo", O_RDONLY | O_LARGEFILE);
```

Writing a file:

We first open a file using open then we use the write() call to actually write our data to a file.

Non-sequential read and write

```
off_t lseek(int fildes, off_t offset, int whence);
```

this simply changes a variable in the OS kernel to tell where to start the reading file from (this is not same as the seek done by a hard disk.

Writing immediately to files

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);
int rc = write(fd, buffer, size);
assert (rc == size);
rc = fsync(fd);
assert(rc == 0);
```

In this case fsync forces all bytes to be written by the write call before moving on.

Renaming Files:

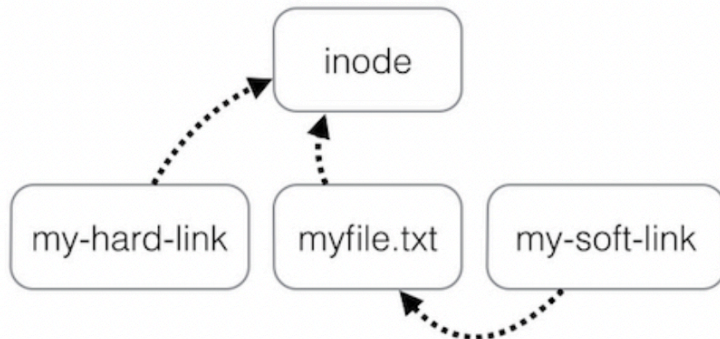
We can use the mv command which then calls the rename function.

```
int fd = open("foo.txt.tmp", O_CREAT | O_WRONLY | O_TRUNC);
write(fd, buffer, size);
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");
```

The rename operation is atomic, so we don't lose a file during a system crash

Metadata

```
struct stat {
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t     st_mode;     /* protection */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t  st_blksize;  /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;   /* number of blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

Inode	An Inode is a persistent data structure kept by the file system that has information about all each of the files.
Removing files	<code>unlink("foo")</code> there is a reference count (aka link count) that will count the number of hard links to a file. Only when the reference count is zero does the OS free the data and inode at the location.
Making/Delete a Directory	<code>mkdir("foo"), rmdir()</code>
Reading Directory:	<pre> opendir(), readdir(), closedir() DIR *dp = opendir("."); struct dirent *d; while ((d = readdir(dp)) != NULL) printf("%d %s\n", (int) d->d_ino, d->d_name) closedir(dp) </pre>
Information stored by each directory (dirent)	<pre> struct dirent { char d_name[256]; /* filename */ ino_t d_ino; /* inode number */ off_t d_off; /* offset to the next dirent */ unsigned short d_reclen; /* length of this record */ unsigned char d_type; /* type of file */ }; </pre>
Symbolic and Hard Links	 <p>If a file gets deleted and there is a symbolic link to it, it becomes a dangling reference. Symlinks can “point” to a directory, hard links cannot go to a directory</p>
Mounting a file system	<p>The <code>mkfs</code> command creates and writes an empty file system starting with a root directory onto that specific disk partition</p> <p>Mount: it takes an existing directory as a target mount point, and pastes a new file system onto the directory tree at that point.</p>